

Lab 7

Morgen Kaufman

Math 241, Week 9

```
# Put all necessary libraries herey
library(tidyverse)
library(tidytext)

# Ensure the textdata package is installed
if (!requireNamespace("textdata", quietly = TRUE)) {
  install.packages("textdata")
}
# Load the textdata package
library(textdata)

# Before knitting your document one last time, you will have to download the AFINN lexicon explicitly
lexicon_afinn()
```

```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon     -2
## 2 abandoned  -2
## 3 abandons   -2
## 4 abducted   -2
## 5 abduction  -2
## 6 abductions -2
## 7 abhor      -3
## 8 abhorred   -3
## 9 abhorrent  -3
## 10 abhors    -3
## # i 2,467 more rows
```

```
lexicon_nrc()
```

```
## # A tibble: 13,872 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
```

```
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # i 13,862 more rows
```

```
library(dplyr)
library(stringr)
```

Due: Friday, March 29th at 5:30pm

Goals of this lab

1. Practice matching patterns with regular expressions.
2. Practice manipulating strings with `stringr`.
3. Practice tokenizing text with `tidytext`.
4. Practice looking at word frequencies.
5. Practice conducting sentiment analysis.

Problem 1: What's in a Name? (You'd Be Surprised!)

1. Load the `babynames` dataset, which contains yearly information on the frequency of baby names by sex and is provided by the US Social Security Administration. It includes all names with at least 5 uses per year per sex. In this problem, we are going to practice pattern matching!

```
library(babynames)
data("babynames")
#?babynames
```

- a. For 2000, find the ten most popular female baby names that start with the letter Z.

```
#Hint: Use
t1 <- babynames %>%
  filter(year == "2000",
         sex == "F",
         str_detect(name, "Z")) %>%
  top_n(10)

t1
```

```
## # A tibble: 10 x 5
##   year sex  name      n    prop
##   <dbl> <chr> <chr> <int>  <dbl>
## 1  2000 F    Zoe    3785 0.00190
## 2  2000 F   Zoey    691 0.000346
## 3  2000 F   Zaria    568 0.000285
## 4  2000 F   Zoie    320 0.000160
## 5  2000 F  Zariah    168 0.0000842
## 6  2000 F   Zion    156 0.0000782
## 7  2000 F  Zainab    142 0.0000712
## 8  2000 F   Zara    121 0.0000607
## 9  2000 F  Zahra    113 0.0000566
## 10 2000 F  Zaira    103 0.0000516
```

b. For 2000, find the ten most popular female baby names that contain the letter z.

```
t2 <- babynames %>%
  filter(year == "2000",
         sex == "F",
         str_detect(name, "[Zz]")) %>%
  top_n(10)

t2
```

```
## # A tibble: 10 x 5
##   year sex   name      n    prop
##   <dbl> <chr> <chr>   <int>  <dbl>
## 1  2000 F   Elizabeth 15094 0.00757
## 2  2000 F   Mackenzie 6348 0.00318
## 3  2000 F    Zoe      3785 0.00190
## 4  2000 F   Mckenzie 2526 0.00127
## 5  2000 F   Makenzie 1613 0.000809
## 6  2000 F   Jazmin    1391 0.000697
## 7  2000 F   Jazmine   1353 0.000678
## 8  2000 F   Lizbeth   817 0.000410
## 9  2000 F   Eliza     759 0.000380
## 10 2000 F   Litzy     722 0.000362
```

c. For 2000, find the ten most popular female baby names that end in the letter z.

```
t3 <- babynames %>%
  filter(year == "2000",
         sex == "F",
         str_detect(name, "z$")) %>%
  top_n(10)

t3
```

```
## # A tibble: 11 x 5
##   year sex   name      n    prop
##   <dbl> <chr> <chr>   <int>  <dbl>
## 1  2000 F    Luz      489 0.000245
## 2  2000 F   Beatriz  357 0.000179
## 3  2000 F  Mercedes 141 0.0000707
## 4  2000 F  Maricruz  96 0.0000481
## 5  2000 F    Liz      72 0.0000361
## 6  2000 F   Inez     69 0.0000346
## 7  2000 F  Odaliz   24 0.0000120
## 8  2000 F  Marycruz  23 0.0000115
## 9  2000 F   Cruz     19 0.00000952
## 10 2000 F   Deniz    16 0.00000802
## 11 2000 F   Taiz     16 0.00000802
```

d. Between your three tables in 1.a - 1.c, do any of the names show up on more than one list? If so, which ones? (Yes, I know you could do this visually but use some joins!)

```
names_1_2 <- inner_join(t1, t2, by = "name")
names_1_3 <- inner_join(t1, t3, by = "name")
names_2_3 <- inner_join(t2, t3, by = "name")
name_z_all <- bind_rows(names_1_2, names_1_3, names_2_3)
name_z_all
```

```
## # A tibble: 1 x 9
##   year.x sex.x name      n.x prop.x year.y sex.y      n.y prop.y
##   <dbl> <chr> <chr> <int>   <dbl> <dbl> <chr> <int>   <dbl>
## 1   2000 F     Zoe    3785 0.00190 2000 F     3785 0.00190
```

e. Verify that none of the baby names contain a numeric (0-9) in them.

```
numeric_names <- babynames %>%
  filter(str_detect(name, "[0-9]"))

numeric_names
```

```
## # A tibble: 0 x 5
## # i 5 variables: year <dbl>, sex <chr>, name <chr>, n <int>, prop <dbl>
```

f. While none of the names contain 0-9, that doesn't mean they don't contain "one", "two", ..., or "nine". Create a table that provides the number of times a baby's name contained the word "zero", the word "one", ... the word "nine".

Notes:

- I recommend first converting all the names to lower case.
- If none of the baby's names contain the written number, there you can leave the number out of the table.
- Use `str_extract()`, not `str_extract_all()`. (We will ignore names where more than one of the words exists.)

Hint: You will have two steps that require pattern matching: 1. Subset your table to only include the rows with the desired words. 2. Add a column that contains the desired word.

```
numberrames <- babynames %>%
  mutate(name = tolower(name),
         zero = str_extract(name, "zero"),
         one = str_extract(name, "one"),
         two = str_extract(name, "two"),
         three = str_extract(name, "three"),
         four = str_extract(name, "four"),
         five = str_extract(name, "five"),
         six = str_extract(name, "six"),
         seven = str_extract(name, "seven"),
         eight = str_extract(name, "eight"),
         nine = str_extract(name, "nine")
  )

numberrames_sub <- numberrames %>%
```

```

filter(rowSums(!is.na(select(., zero, one, two, three, four, five, six, seven, eight, nine))) > 0)

numberrnames_count <- numberrnames_sub %>%
  summarise(
    zero = sum(!is.na(zero)),
    one = sum(!is.na(one)),
    two = sum(!is.na(two)),
    three = sum(!is.na(three)),
    four = sum(!is.na(four)),
    five = sum(!is.na(five)),
    six = sum(!is.na(six)),
    seven = sum(!is.na(seven)),
    eight = sum(!is.na(eight)),
    nine = sum(!is.na(nine))
  )

numberrnames_count

```

```

## # A tibble: 1 x 10
##   zero one two three four five six seven eight nine
##   <int> <int> <int> <int> <int> <int> <int> <int> <int> <int>
## 1     4 10210  288   58    2    0  106   50  356  807

```

g. Which written number or numbers don't show up in any of the baby names?

Five does not show up at all.

h. Create a table that contains the names and their frequencies for the two least common written numbers.

```

least_common_numbers_count <- numberrnames %>%
  filter(rowSums(!is.na(select(., zero, four))) > 0) %>%
  count(name)

least_common_numbers_count

```

```

## # A tibble: 4 x 2
##   name      n
##   <chr>  <int>
## 1 balfour    2
## 2 luzero    2
## 3 zero      1
## 4 zeron     1

```

```

least_common_numbers <- numberrnames %>%
  filter(rowSums(!is.na(select(., zero, four))) > 0)

least_common_numbers

```

```

## # A tibble: 6 x 15
##   year sex name      n      prop zero one two three four five six
##   <dbl> <chr> <chr>  <int>  <dbl> <chr> <chr> <chr> <chr> <chr> <chr> <chr>

```

```
## 1 1914 M      balfour      5 0.00000732 <NA> <NA> <NA> <NA> four <NA> <NA>
## 2 1928 M      balfour      5 0.00000438 <NA> <NA> <NA> <NA> four <NA> <NA>
## 3 1990 F      luzero       6 0.00000292 zero <NA> <NA> <NA> <NA> <NA> <NA>
## 4 1991 F      luzero       5 0.00000246 zero <NA> <NA> <NA> <NA> <NA> <NA>
## 5 2007 M      zeron        6 0.00000271 zero <NA> <NA> <NA> <NA> <NA> <NA>
## 6 2017 M      zero         7 0.00000357 zero <NA> <NA> <NA> <NA> <NA> <NA>
## # i 3 more variables: seven <chr>, eight <chr>, nine <chr>
```

- i. List out the names that contain no vowels (consider “y” to be a vowel).

```
# Filter names containing no vowels
names_no_vowels <- babynames %>%
  filter(!str_detect(name, "[aeiouyAEIOUY]")) %>%
  count(name)

names_no_vowels
```

```
## # A tibble: 43 x 2
##   name      n
##   <chr> <int>
## 1 Bb         5
## 2 Bg         6
## 3 Bj        49
## 4 Cj        62
## 5 Dj        49
## 6 Jb        69
## 7 Jc        97
## 8 Jd        85
## 9 Jj        44
## 10 Jl         6
## # i 33 more rows
```

Problem 2: Tidying the “Call of the Wild”

Did you read “Call of the Wild” by Jack London? If not, [read the first paragraph of its wiki page](#) for a quick summary and then let’s do some text analysis on this classic! The following code will pull the book into R using the `gutenbergr` package.

```
library(gutenbergr)
wild <- gutenberg_download(215)
```

- a. Create a tidy text dataset where you tokenize by words.

```
wild_token <- wild %>%
  unnest_tokens(output = word, input = text)
```

- b. Find the frequency of the 20 most common words. First, remove stop words.

```
data("stop_words")

wild20 <- wild_token %>%
  anti_join(stop_words) %>%
  group_by(word) %>%
  summarise(frequency = n()) %>%
  arrange(desc(frequency)) %>%
  top_n(20)

wild20
```


```
## # A tibble: 20 x 2
##   word      frequency
##   <chr>      <int>
## 1 buck        313
## 2 dogs        118
## 3 thornton     81
## 4 dog          79
## 5 time         77
## 6 day          70
## 7 life         64
## 8 sled         63
## 9 half         60
## 10 spitz        60
## 11 head         54
## 12 camp         53
## 13 françois     51
## 14 feet         49
## 15 buck's       47
## 16 trail        42
## 17 days         41
## 18 eyes         40
## 19 john         40
## 20 night        40
```

c. Create a bar graph and a word cloud of the frequencies of the 20 most common words.

```
library(wordcloud)

ggplot(data = wild20, aes(x = reorder(word, frequency), y = frequency)) +
  geom_bar(stat = "identity", fill = "#e0218a") +
  labs(x = "Word", y = "Frequency") +
  coord_flip() +
  ggtitle("Top 20 Most Common Words") +
  theme_minimal()
```

Top 20 Most Common Words



```
wordcloud(words = wild20$word,  
  freq = wild20$frequency,  
  scale = c(4, 1),  
  rot.per = .5,  
  colors = "#e0218a",  
  random.order = FALSE)
```




- d. Explore the sentiment of the text using three of the sentiment lexicons in `tidytext`. What does your analysis say about the sentiment of the text?

Notes:

- Make sure to NOT remove stop words this time.
- `afinn` is a numeric score and should be handled differently than the categorical scores.

Afinn

```
wild_group <- wild_token %>%
  group_by(word) %>%
  summarise(frequency = n())

afinn_wild <- wild_group %>%
  left_join(get_sentiments("afinn")) %>%
  filter(!is.na(value)) %>%
  arrange(desc(value)) %>%
  group_by(word)
afinn_wild
```

```
## # A tibble: 528 x 3
## # Groups:   word [528]
```

```
##      word      frequency value
##      <chr>         <int> <dbl>
## 1 miracle          2      4
## 2 terrific          1      4
## 3 triumph          1      4
## 4 win              1      4
## 5 wonderful        1      4
## 6 adore            1      3
## 7 affection        2      3
## 8 beautiful        4      3
## 9 best             10      3
## 10 cheery          1      3
## # i 518 more rows
```

```
afinn_wild %>%
  group_by(value) %>%
  summarise(n())
```

```
## # A tibble: 9 x 2
##   value 'n()'
##   <dbl> <int>
## 1    -5      1
## 2    -4      1
## 3    -3     56
## 4    -2    175
## 5    -1     80
## 6     1     71
## 7     2    101
## 8     3     38
## 9     4      5
```

```
afinn_wild %>%
  ungroup() %>%
  summarise(
    mean = mean(value))
```

```
## # A tibble: 1 x 1
##   mean
##   <dbl>
## 1 -0.379
```

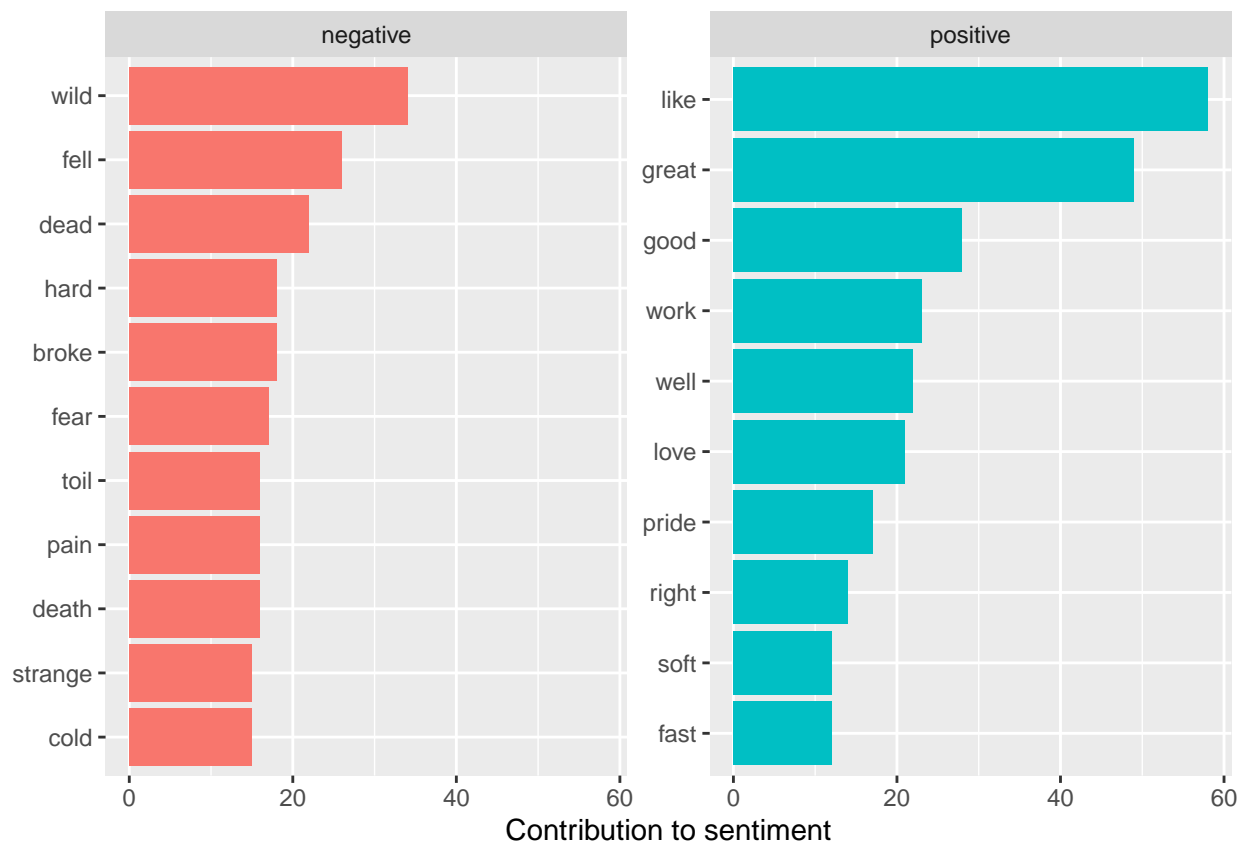
Bing

```
bing_wild <- wild_token %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  ungroup()
bing_wild
```

```
## # A tibble: 892 x 3
##   word sentiment      n
##   <chr> <chr>    <int>
```

```
## 1 like positive 58
## 2 great positive 49
## 3 wild negative 34
## 4 good positive 28
## 5 fell negative 26
## 6 work positive 23
## 7 dead negative 22
## 8 well positive 22
## 9 love positive 21
## 10 broke negative 18
## # i 882 more rows
```

```
bing_wild %>%
  group_by(sentiment) %>%
  slice_max(n, n = 10) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  labs(x = "Contribution to sentiment",
       y = NULL)
```

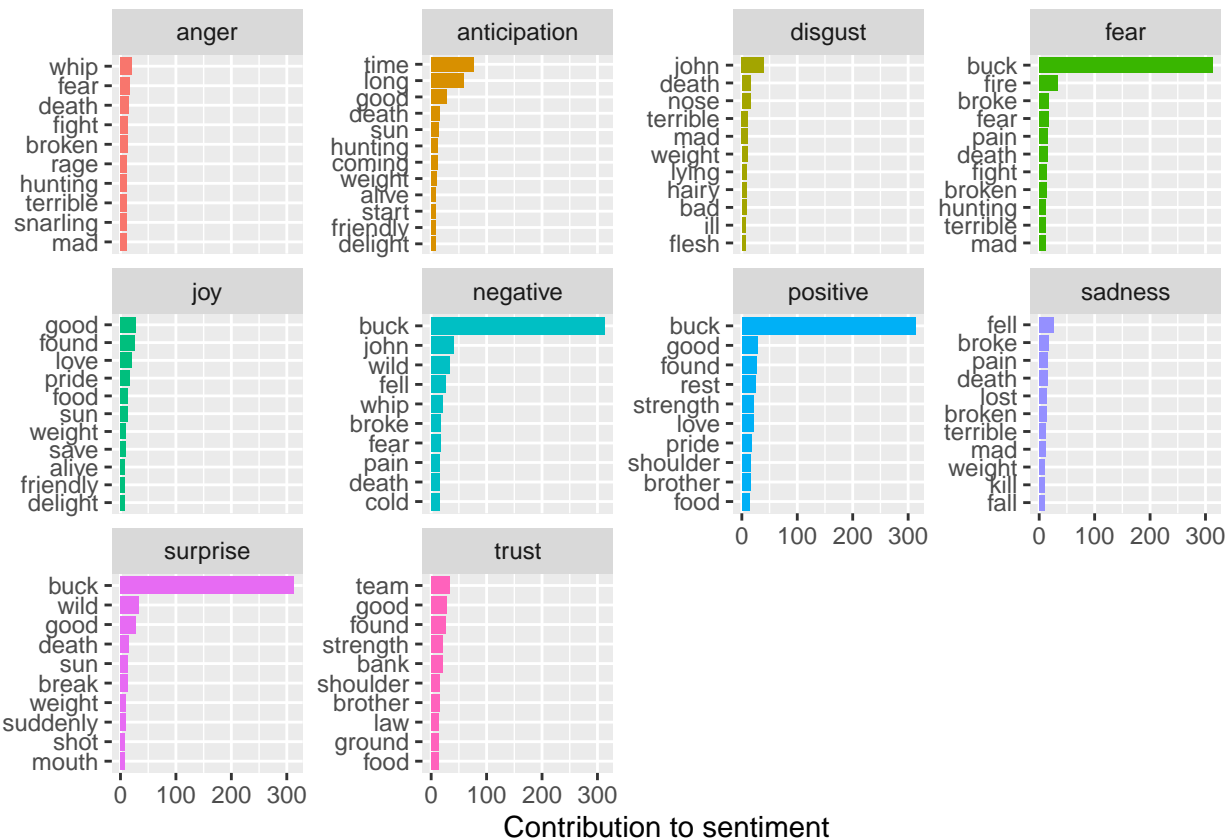


NRC

```
nrc_wild <- wild_token %>%
  inner_join(get_sentiments("nrc")) %>%
  count(word, sentiment) %>%
  ungroup()
nrc_wild
```

```
## # A tibble: 2,613 x 3
##   word      sentiment      n
##   <chr>      <chr>      <int>
## 1 abandonment anger          1
## 2 abandonment fear           1
## 3 abandonment negative        1
## 4 abandonment sadness         1
## 5 abandonment surprise        1
## 6 ability     positive         2
## 7 absent      negative         1
## 8 absent      sadness          1
## 9 abundance   anticipation         2
## 10 abundance  disgust           2
## # i 2,603 more rows
```

```
nrc_wild %>%
  group_by(sentiment) %>%
  slice_max(n, n = 10) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  labs(x = "Contribution to sentiment",
       y = NULL)
```



The general sentiment seems to be quite neutral. While there looks like a slight negative skew with `afinn` and a slight positive skew with `bing` in general the sentiment is relatively neutral with a good balance of negative and positive words.

- e. If you didn't do so in 2.d, compute the average sentiment score of the text using `afinn`. Which positive words had the biggest impact? Which negative words had the biggest impact?

```
wild_token %>%
  group_by(word) %>%
  summarise(frequency = n()) %>%
  left_join(get_sentiments("afinn")) %>%
  filter(!is.na(value)) %>%
  arrange(desc(frequency)) %>%
  top_n(5, wt = frequency)
```

```
## # A tibble: 5 x 3
##   word frequency value
##   <chr>      <int> <dbl>
## 1 no          95    -1
## 2 like        58     2
## 3 great       49     3
## 4 fire        33    -2
## 5 good        28     3
```

Top Positive: like, great, good

Top Negative: no, fire

- f. You should have found that “no” was an important negative word in the sentiment score. To know if that really makes sense, let’s turn to the raw lines of text for context. Pull out all of the lines that have the word “no” in them. Make sure to not pull out extraneous lines (e.g., a line with the word “now”).

```
nowild <- wild %>%
  filter(str_detect(text, "\\bno\\b"))
print(nowild)
```

```
## # A tibble: 83 x 2
##   gutenbergs_id text
##   <int> <chr>
## 1      215 solitary man, no one saw them arrive at the little flag station~
## 2      215 that it was the club, but his madness knew no caution. A dozen ~
## 3      215 "He's no slouch at dog-breakin', that's wot I say," one of the ~
## 4      215 all, that he stood no chance against a man with a club. He had ~
## 5      215 in the red sweater. "And seem' it's government money, you ain't~
## 6      215 animal. The Canadian Government would be no loser, nor would its
## 7      215 while he developed no affection for them, he none the less grew
## 8      215 The other dog made no advances, nor received any; also, he did ~
## 9      215 knew no law but the law of club and fang.
## 10     215 full-grown wolf, though not half so large as she. There was no ~
## # i 73 more rows
```

- g. Draw some conclusions about how “no” is used in the text.

No within this text seems to be used as a modifier for adjectives and verbs. In some cases when it occurs while it does have a slight negative connotation in others it is actually positive such as “no more trouble” and “no matter what the odds.” As such it seems that no while greatly contributing to the calculated sentiment is quite neutral in most cases and works purely as a modifier.

- h. We can also look at how the sentiment of the text changes as the text progresses. Below, I have added two columns to the original dataset. Now I want you to do the following wrangling:

- Tidy the data (but don’t drop stop words).
- Add the word sentiments using `bing`.
- Count the frequency of sentiments by index.
- Reshape the data to be wide with the count of the negative sentiments in one column and the positive in another, along with a column for index.
- Compute a sentiment column by subtracting the negative score from the positive.

```
wild_time <- wild %>%
  mutate(line = row_number(), index = floor(line/45) + 1)

wild_token <- wild_time %>%
  unnest_tokens(output = word, input = text)


wild_sent <- wild_token %>%
  inner_join(get_sentiments("bing"))
wild_sent
```

```
## # A tibble: 2,351 x 5
##   gutenber_id line index word      sentiment
##         <int> <int> <dbl> <chr>    <chr>
## 1         215     6     1 wild      negative
## 2         215    15     1 primitive negative
## 3         215    18     1 won        positive
## 4         215    19     1 toil       negative
## 5         215    20     1 love       positive
## 6         215    26     1 primitive negative
## 7         215    32     1 strain    negative
## 8         215    35     1 trouble   negative
## 9         215    37     1 strong    positive
## 10        215    37     1 warm      positive
## # i 2,341 more rows
```

```
wild_score <- wild_sent %>%
  group_by(index) %>%
  summarise(count = n(), total_pos = sum(sentiment == "positive"), total_neg = sum(sentiment == "negative"),
    mutate(score = total_pos - total_neg)
```

i. Create a plot of the sentiment scores as the text progresses.

```
wild_score %>%
  ggplot(aes(x = index, y = score)) +
  geom_line() +
  geom_smooth()
```



- j. The choice of 45 lines per chunk was pretty arbitrary. Try modifying the index value a few times and recreating the plot in i. Based on your plots, what can you conclude about the sentiment of the novel as it progresses?

```
wild_time <- wild %>%
  mutate(line = row_number(), index = floor(line/10) + 1)

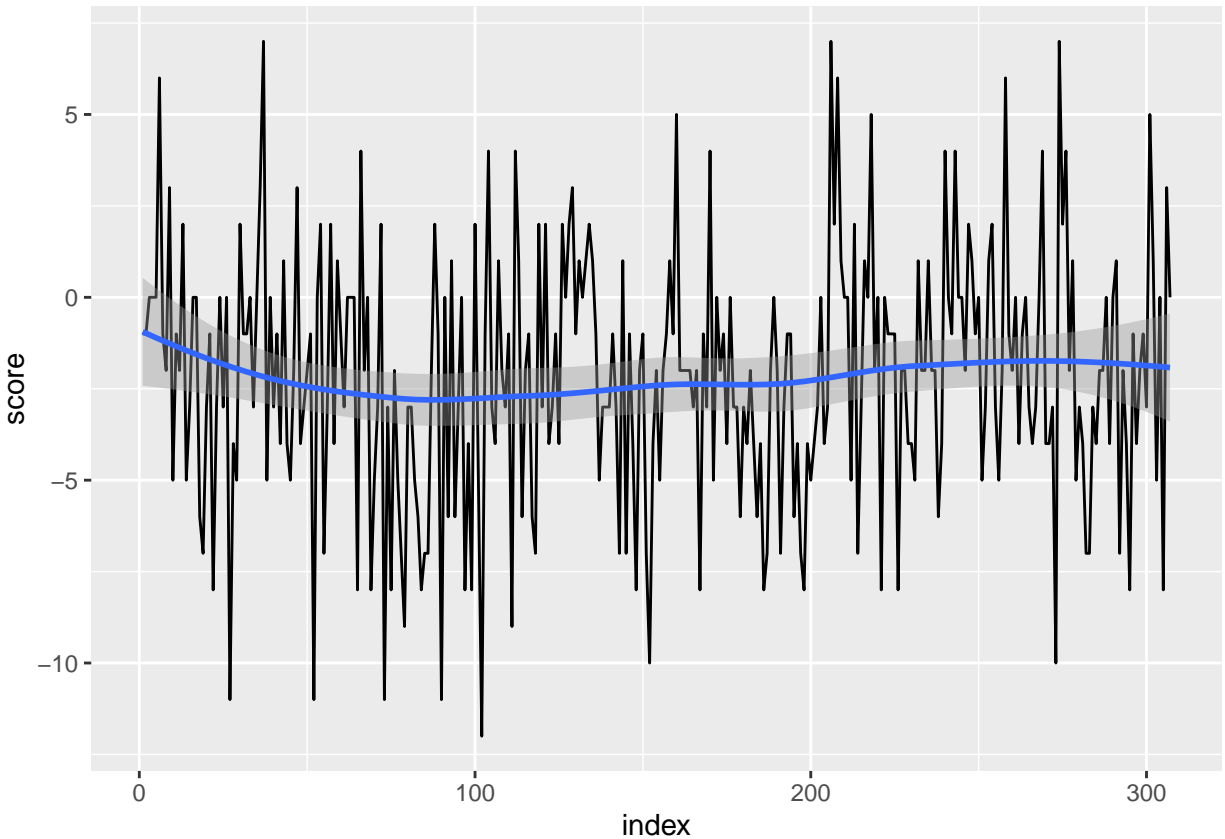
wild_token <- wild_time %>%
  unnest_tokens(output = word, input = text)

wild_sent <- wild_token %>%
  inner_join(get_sentiments("bing"))
wild_sent
```

```
## # A tibble: 2,351 x 5
##   gutenber_id line index word      sentiment
##         <int> <int> <dbl> <chr>    <chr>
## 1         215     6     1 wild      negative
## 2         215    15     2 primitive negative
## 3         215    18     2 won        positive
## 4         215    19     2 toil       negative
## 5         215    20     3 love       positive
## 6         215    26     3 primitive negative
## 7         215    32     4 strain    negative
## 8         215    35     4 trouble   negative
## 9         215    37     4 strong    positive
## 10        215    37     4 warm      positive
## # i 2,341 more rows
```

```
wild_score <- wild_sent %>%
  group_by(index) %>%
  summarise(count = n(), total_pos = sum(sentiment == "positive"), total_neg = sum(sentiment == "negative"))
mutate(score = total_pos - total_neg)

wild_score %>%
  ggplot(aes(x = index, y = score)) +
  geom_line() +
  geom_smooth()
```

Throughout the novel there is a constant and drastic shift from positive to negative in scores. However, in the case of negative, the lowest score is much more than the highest positive score. Additionally while there are near constant fluctuations, when plotted and averaged out the score is fairly neutral.

k. Let's look at the bigrams (2 consecutive words). Tokenize the text by bigrams.

```
wild_bigrams <- wild %>%
  unnest_tokens(output = bigram, input = text, token = "ngrams", n = 2) %>%
  na.omit()
```

```
wild_bigrams
```

```
## # A tibble: 29,442 x 2
##   gutenber_id bigram
##   <int> <chr>
## 1      215 the call
## 2      215 call of
## 3      215 of the
## 4      215 the wild
## 5      215 by jack
## 6      215 jack london
## 7      215 chapter i
## 8      215 i into
## 9      215 into the
## 10     215 the primitive
## # i 29,432 more rows
```

1. Produce a sorted table that counts the frequency of each bigram and notice that stop words are still an issue.

```
wild_bigrams %>%  
  group_by(bigram) %>%  
  summarise(frequency = n()) %>%  
  arrange(desc(frequency))
```

```
## # A tibble: 18,907 x 2  
##   bigram    frequency  
##   <chr>      <int>  
## 1 of the      233  
## 2 in the      172  
## 3 he was      127  
## 4 to the      116  
## 5 it was      107  
## 6 and the       95  
## 7 on the       80  
## 8 he had       77  
## 9 at the       68  
## 10 into the    65  
## # i 18,897 more rows
```